



CIS  
1905

# Lecture 6: Smart Pointers

Beyond simple ownership



# Logistics

HW3 deadline changes

Final Project Proposals

[pollev.com/cis1905rust915](https://pollev.com/cis1905rust915)



# Some Reminders

## Ownership Rules

1. Each value has an owner
2. There can only be one owner at a time
3. When the owner goes out of scope, the value is dropped

## Borrowing Rules

1. At any given time, you can have either one mutable reference or many immutable references
2. References must always be valid

# Some Reminders

## Ownership Rules

1. Each value has an owner
- ~~2. There can only be one owner at a time~~
3. When the owner goes out of scope, the value is dropped

## Borrowing Rules

- ~~1. At any given time, you can have either one mutable reference or many immutable references~~
2. References must always be valid

We can break\* some of these rules with Smart Pointers.

The other ones require `unsafe` rust.

# Smart Pointers

`Box<T>`

**Single ownership on  
the heap**

`Cow<'a, B>`

**Lazy copies**

`Rc<T> / Arc<T>`

**Multiple ownership  
(without mutation)**

`Vec<T> / ...`

**Dynamically sized  
collections**

`Rc<RefCell<T>> / ...`

**Multiple ownership  
(with mutation)**

`&[T], &dyn T`


**Wide pointers**

# Box<T>

- Just allocates T on the heap
- T is owned by the Box; when the Box is dropped, T's heap space is freed
- The Box's control of its heap space is exclusive
- (if you're familiar with C++, this is like `std::unique_ptr`)

**Enforces all of Rust's borrowing rules at compile time**

```
fn main() {  
    // Allocate the value 10.0 on the heap  
    let x = Box::new(10.0);  
    assert_eq!(*x, 10.0);  
  
    // Deallocate the heap value  
    drop(x);  
}
```



## Box<T> Use Case: Recursive Data Types

```
struct BTree<T> {  
    value: T,  
    left_child: Option<Self>,  
    right_child: Option<Self>,  
}
```

💡 recursive type `BTree` has infinite size

💡 recursive without indirection

$$\text{size}(\text{BTree}) = \text{size}(T) + \text{size}(\text{Option}<\text{BTree}>) = \infty$$

Fix the infinite recursion with **Box**:

```
struct BTree<T> {  
    value: T,  
    left_child: Option<Box<Self>>,  
    right_child: Option<Box<Self>>,  
}
```

$$\begin{aligned}\text{size}(\text{BTree}) &= \text{size}(T) + \text{size}(\text{Box}<\text{BTree}>) \\ &= \text{size}(\text{BTree}) = \text{size}(T) + \text{size}(\text{usize})\end{aligned}$$

# Rc<T>

What if you want more than one owner?

**Rc** stands for “reference counted”.

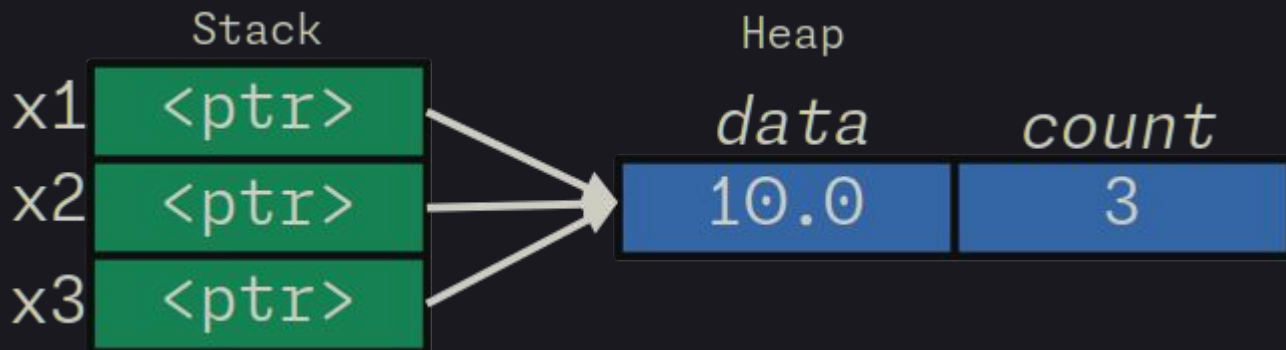
- The data lives on the **heap**
- The references live on the **stack**

Like a normal reference, but *everyone* is an owner—and you only get immutable access to the data



# How does it work? (simplified version)

- `Rc::new(data)`
  - Allocates space on the heap for `data`, sets count to 1
- `impl Clone`
  - Copies the *pointer* and increments `count` by 1
  - Doesn't actually clone the `data`!
- `impl Drop`
  - Decrement `count`
  - If reference count is 0, then free the `data`



# Doesn't this break the ownership rules?

Kind of, but not really

- Multiple owners point to the same data, but it is **immutable**
  - In this way, **Rc** behaves like a reference
- If all the owners go out of scope, the data is dropped
  - In this way, **Rc** behaves like an owned value

So yes, it breaks the rules, but it also maintains safety.

```
fn main() {  
    let a = Rc::new(5.0);  
    // data gets allocated  
    let b = a.clone();  
    {  
        let c = b.clone();  
        println!("{}", c);  
    }  
    drop(b);  
    drop(a);  
    // data gets dropped  
}
```

# What does this look like in practice?

You should use **Rc** when:

- You have some *really big* piece of data
- It's expensive to **clone**
- Lots of things need **access** to it
- It doesn't make sense to have a **single** owner
- You don't know **how long** it will be needed
- You want it to be **easy**

*Cloning is free!* →

```
/// A really big texture that's expensive to clone around
pub struct Texture;
impl Texture {
    pub fn load(path: &str) -> Rc<Self> {
        unimplemented!()
    }
}

/// A sprite in your game
pub struct Sprite {
    pub name: String,
    pub texture: Rc<Texture>,
}
impl Sprite {
    pub fn new(name: impl ToString, texture: Rc<Texture>) -> Self {
        Self {
            name: name.to_string(),
            texture,
        }
    }
}

fn main() {
    let texture = Texture::load("big_texture.png");

    let player = Sprite::new("Main Player", texture.clone());

    let enemy1 = Sprite::new("Enemy", texture.clone());
    let enemy2 = Sprite::new("Enemy", texture.clone());
    let enemy3 = Sprite::new("Enemy", texture.clone());
}
```

# Cyclical references

Beware of cyclical references!

The reference count will never hit zero, so your data will never be freed.

- This is a memory leak

Solution: **Weak**<T>

- Same as a **Rc**<T>, but doesn't increment the count (the data is **not** guaranteed to be there)



```
let weak = Rc::downgrade(&strong);
```

# Detour: Cell<T>

- Cell is not a pointer type (but it builds towards one...)
- Let's you *mutate* an *immutable* value in a controlled way

In general, it's useless...

- You can't read the value, only replace it with a new one

But for **Copy** types...

- You can get the value by copying it, so it becomes useful

Doesn't this break the rules?

- Single owner: not broken, cells are still owned values
- One writer XOR multiple readers: not broken, since you're never actually mutating anything; you're just *replacing* one value with another

## Detour: RefCell<T>

- Also not a pointer! Allows us to *mutate* any *immutable* value, not just **Copy** ones!
- Enforces the borrow checker's rules *at runtime*

```
let test: RefCell<String> = RefCell::new(String::from_str("Hello").unwrap());
{
    let inner_value: Ref<'_, String> = test.borrow();
    println!("{inner_value}") // Prints "Hello"
}
{
    let mut mutable_ref: RefMut<'_, String> = test.borrow_mut();
    mutable_ref.push_str(string: " World!");
}
{
    let inner_value: Ref<'_, String> = test.borrow();
    println!("{inner_value}") // Prints "Hello World!", even though test is immutable!
}
```

# Detour: RefCell<T>

So this is fine:

not mut

```
let test: RefCell<String> = RefCell::new(String::from_str("Hello").unwrap());
{
    let inner_value: Ref<'_, String> = test.borrow();
    println!("{inner_value}") // Prints "Hello"
}
{
    let mut mutable_ref: RefMut<'_, String> = test.borrow_mut();
    mutable_ref.push_str(string: " World!");
}
{
    let inner_value: Ref<'_, String> = test.borrow();
    println!("{inner_value}") // Prints "Hello World!", even though test is immutable!
}
```

## Detour: RefCell<T>

But this isn't:

```
{  
    let immutable_ref: Ref<'_, String> = test.borrow();  
    let mut mutable_ref: RefMut<'_, String> = test.borrow_mut(); // This will panic!!  
    println!("{immutable_ref}");  
    mutable_ref.push_str(string: "oops");  
}
```

*Why isn't this ok?*



# Rc<RefCell<T>>

*What if you want more than one owner and mutability via runtime borrow-checking?*

- For when you really want to wing it 😊
- ~Almost~ a Java reference; like `std::shared_ptr` in C++
- Allows multiple owners to mutate the same heap location, so long the borrow rules are upheld at runtime

# Let's make a graph!

*Attempt 1*

```
struct Node {  
    value: i32,  
    children: Vec<Box<Node>>,  
}
```

Well, that's not right...

- Each node needs to own its children, so we can't make cycles :(
- We need *multiple ownership*!

# Let's make a graph!

Attempt 2

```
struct Node {  
    value: i32,  
    children: Vec<Rc<Node>>,  
}
```

Well, that's not right...

- We can't modify anything anymore! Our references to nodes are immutable
- So we still can't make a cycle...
- We need *runtime-checked borrowing*!

# Let's make a graph!

*Attempt 3*

```
struct Node {  
    value: i32,  
    children: Vec<Rc<RefCell<Node>>>,  
}
```

Yay cycles! (But watch out for cycles)

```
let node1: Rc<RefCell<Node>> = Rc::new(RefCell::new(Node { value: 1, children: vec![] }));  
let node2: Rc<RefCell<Node>> = Rc::new(RefCell::new(Node { value: 2, children: vec![] }));  
RefCell::borrow_mut(&node1).children.push(Rc::clone(&node2));  
RefCell::borrow_mut(&node2).children.push(node1);
```



Cow< ' a, B>  
what?



# Clone On Write

\*Essentially a “lazy clone”

\*You probably won't use this one

```
pub enum Cow<'a, B> { ... }
```

## Borrowed(&'a B)

When you call `Cow::from(x)`, you get a `Borrowed(&x)`.

This is essentially free since all you did is take a reference. It will stay this way until you need to mutate the data.

## Owned(<B as ToOwned>::Owned)

There are two ways to get this

1. `Cow::to_mut(&mut self)`  
Clones the data, and gives you a mutable reference
2. `Cow::into_owned(self)`  
Clones the data and gives it to you

`Cow` only clones the data once you need mutable access, then it stays that way

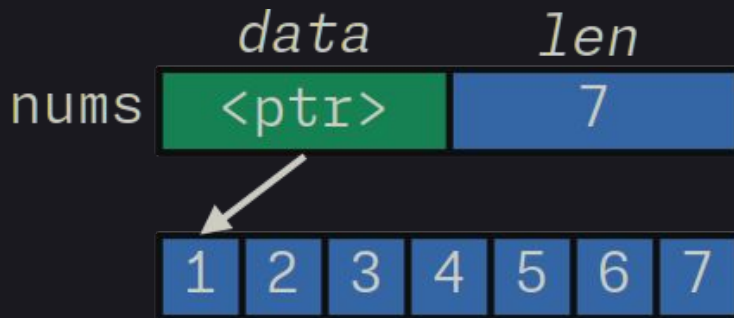
# Wide pointers

Some pointers are called “wide pointers”, which means they carry extra information

There are two kinds of wide pointers

- Slices (`&[T]`), which store the **length** of the slice (this also includes `&str`)
- Trait objects (`&dyn T`), which store a pointer to the **vtable**

The pointers are “wide” because they take up 2 words instead of just 1.



Slice example

# Wide pointers

They don't have to be behind *references*, they just have to be *pointers*

Some weird looking but syntactically-correct examples:

```
Box<dyn Error>
```

Owning pointer to some type that implements Error

```
Rc<str>
```

Reference-counted pointer to an immutable string slice

```
Cow<[i32]>
```

Clone-on-writeable integer array slice



# Thread-Safe Smart Pointers

Most of the normal smart pointers are not thread safe for efficiency reasons, so Rust gives us thread-safe alternatives

## Normal

`Rc`

`Cell<T>`

`RefCell<T>`

## Thread-Safe

`Arc`: atomically reference counted `Rc`

`AtomicT`: (like `AtomicI32`)

`RWLock<T>`: one writer XOR many readers

`Mutex<T>`: one writer XOR one reader

# unsafe

If smart pointers break the ownership and borrowing rules, how are they implemented?

1. Unsafe rust
  - a. Lets you do bad things (i.e. work with raw (C-style) pointers, work with uninitialized memory, transmute types, etc.)
  - b. Must be contained within an `unsafe` block
  - c. By writing `unsafe` code, you “promise” not to cause undefined behavior or otherwise break Rust’s contracts
2. Compiler intrinsics
  - a. Some types (such as `Box`) are so fundamental that they are implemented at the compiler level

```
fn main() {  
    let x: Box<i32> = unsafe { Box::new(MaybeUninit::uninit().assume_init()) };  
    println!("x = {}", x);  
}
```

`x = -1824071680` (undefined behavior)

# Challenge: Binary Tree


Implement a binary tree on your own!

Each vertex should have associated data and two optional children.

You should support the following operations:

- `new(data: T)`
- `add_left_child(child: Self)`
- `add_right_child(child: Self)`
- `contains(data: T)`

You may want to `#[derive(Debug)]` to test your tree.



## Sample Solution

```
#[derive(Clone, Debug)]  
struct BTree<T> {  
    value: T,  
    left: Option<Box<BTree<T>>>,  
    right: Option<Box<BTree<T>>>,  
}
```

```
impl<T> BTree<T> {
    fn new(value: T) -> Self {
        Self {
            value,
            left: None,
            right: None,
        }
    }

    fn add_left_child(&mut self, child: Self) {
        self.left = Some(Box::new(child));
    }

    fn add_right_child(&mut self, child: Self) {
        self.right = Some(Box::new(child));
    }

    fn contains(&self, value: &T) -> bool
    where
        T: PartialEq,
    {
        return self.value == *value
            || self.left.as_ref().is_some_and(|l| l.contains(value))
            || self.right.as_ref().is_some_and(|r| r.contains(value));
    }
}
```

# Challenge: Improved Binary Tree


Now, make it hold pointers to both the parent *and* the child.

You should support the following additional operations:

- `get_left_child(&self)`
- `get_right_child(&self)`
- `get_parent(&self)`

Hint: change `new` to return type of `Rc<RefCell<Self>>`

Hint: change `add_left_child` and `add_right_child` to take in `parent` and `child` instead of `self` and `child`



## Sample Solution

```
#[derive(Clone, Debug)]
struct BTree<T> {
    value: T,
    left: Option<Rc<RefCell<BTree<T>>>>,
    right: Option<Rc<RefCell<BTree<T>>>>,
    parent: Option<Rc<RefCell<BTree<T>>>>,
}
```

# Sample Solution

```
impl<T> BTree<T> {  
    fn new(value: T) -> Rc<RefCell<Self>> {  
        Rc::new(RefCell::new(Self {  
            value,  
            left: None,  
            right: None,  
            parent: None,  
        })))  
    }  
  
    fn add_left_child(parent: Rc<RefCell<Self>>, child: Rc<RefCell<Self>>) {  
        child.borrow_mut().parent = Some(parent.clone());  
        parent.borrow_mut().left = Some(child);  
    }  
  
    fn add_right_child(parent: Rc<RefCell<Self>>, child: Rc<RefCell<Self>>) {  
        child.borrow_mut().parent = Some(parent.clone());  
        parent.borrow_mut().right = Some(child);  
    }  
}
```



# Sample Solution

```
fn contains(&self, value: &T) -> bool
where
    T: PartialEq,
{
    return self.value == *value
        || self
            .left
            .as_ref()
            .is_some_and(|l| l.borrow().contains(value))
        || self
            .right
            .as_ref()
            .is_some_and(|r| r.borrow().contains(value));
}
```

# Sample Solution

```
fn get_left_child(&self) -> Option<Rc<RefCell<Self>>> {  
    self.left.clone()  
}  
  
fn get_right_child(&self) -> Option<Rc<RefCell<Self>>> {  
    self.right.clone()  
}  
  
fn get_parent(&self) -> Option<Rc<RefCell<Self>>> {  
    self.parent.clone()  
}  
}
```